

Advanced Programming in the UNIX Environment — *Files and Directories*

Hop Lee
hoplee@bupt.edu.cn

SCHOOL OF INFORMATION AND COMMUNICATION ENGINEERING



Table of Contents I

File Status

File Types

Set-User-ID and Set-Group-ID

File Access Permissions

Ownership of New Files and Directories

ACCESS and FACCESSAT Functions

UMASK Function

CHMOD, FCHMOD, and FCHMODAT Functions

Sticky Bit

CHOWN, FCHOWN, FCHIWNAT, and LCHOWN Functions

File Size

File Truncation

Filesystems

LINK, LINKAT, UNLINK, UNLINKAT, and REMOVE Functions



Table of Contents II

RENAME and RENAMEAT Functions

Symbolic Link

Creating and Reading Symbolic Links

File Time

FUTIMENS, UTIMESAT, and UTIMES Functions

MKDIR, MKDIRAT, and RMDIR Functions

Reading Directories

CHDIR, FCHDIR, and GETCWD Functions

Device Special Files

Summary



File Status I

- ▶ The `stat` function returns a structure of information about the given file.

```
1 #include <sys/types.h>
2 #include <sys/stat.h>
3 int stat(const char *restrict pathname,
4          struct stat *restrict buf);
5 int fstat(int filedes, struct stat *buf);
6 int lstat(const char *restrict pathname,
7           struct stat *restrict buf);
8 int fstatat(int fd, const char *restrict pathnam
9             struct stat *restrict buf, int flag)
```



File Status II

- ▶ Given a *pathname*, the *stat* function returns a structure of information about the named file. The *fstat* function obtains information about the file that is already open on the descriptor *fd*. The *lstat* function is similar to *stat*, but when the named file is a symbolic link, *lstat* returns information about the symbolic link itself, not the file referenced by the symbolic link.
- ▶ The *fstatat* function provides a way to return the file statistics for a pathname relative to an open directory represented by the *fd* argument. The *flag* argument controls whether symbolic links are followed.



File Status III

- ▶ The *buf* argument is a pointer to a structure that we must supply. The functions fill in the structure. The definition of the structure can differ among implementations, but it could look like

```
1 struct stat {  
2     mode_t st_mode; /* file type & mode (permissions) */  
3     ino_t st_ino; /* i-node number (serial number) */  
4     dev_t st_dev; /* device number (file system) */  
5     dev_t st_rdev; /* device number for special files */  
6     nlink_t st_nlink; /* number of links */  
7     uid_t st_uid; /* user ID of owner */  
8     gid_t st_gid; /* group ID of owner */  
9     off_t st_size; /* size in bytes, for regular files */  
10    struct timespec st_atim; /* time of LAT */  
11    struct timespec st_mtim; /* time of LMT */  
12    struct timespec st_ctim; /* time of LCT */  
13    blksize_t st_blksize; /* best I/O block size */
```



File Status IV

```
14     blkcnt_t st_blocks; /* no. of disk blocks allocated */  
15 };
```

- ▶ The `timespec` structure type defines time in terms of seconds and nanoseconds. It includes at least the following fields:

```
1 time_t tv_sec;  
2 long tv_nsec;
```

- ▶ We'll go through each member of this structure to examine the attributes of a file in the rest of this chapter.



File Types I

- ▶ There are seven file types in most UNIX system:
 - ▶ Regular file
 - ▶ Directory file
 - ▶ Character device file
 - ▶ Block device file
 - ▶ FIFO
 - ▶ Socket
 - ▶ Symbolic link.
- ▶ Example (Figure 4.3, `filedir/filetype.c`).
- ▶ The type of a file is encoded in the `st_mode` member of the `stat` structure. We can determine the file type by those macros listed in Figure 4.1.



File Types II

- ▶ Next table shows the counts and percentages for a Linux system that is used as a single-user workstation. This data was obtained from the program shown in Section 4.22.

Table: Counts and percentages of different file types

File type	Count	Percentage(%)
regular file	415,803	79.77
directory	62,197	11.93
symbolic link	40,018	8.25
character special	155	0.03
block special	47	0.01
socket	45	0.01
FIFO	0	0.00



Set-User-ID and Set-Group-ID

- ▶ Each process has six or more IDs associated with it: real, effective, and saved set.
- ▶ The real ID identify who we really are.
- ▶ The effective ID and supplementary group ID determine our file access permissions.
- ▶ The saved set ID contain copies of the effective ID when a program is executed. See §8.11.
- ▶ The **set-user-ID** bit and **set-group-ID** bit are contained in the file's `st_mode` value. These two bits can be tested against the constants `S_ISUID` and `S_ISGID`.



File Access Permissions I

- ▶ The `st_mode` value also encodes the access permission bits for the file.
- ▶ Read permission and execute permission for a directory mean different things. Read permission let us can use “`ls`” command, execute permission lets us pass through the directory when it is a component of a pathname that we trying to access.
- ▶ The execute permission bit for a directory is often called the search bit.
- ▶ We must have write permission for a file to specify the `O_TRUNC` flag in the `open` function.



File Access Permissions II

- ▶ To delete an existing file, we need write and execute permission in the directory containing the file, but do **NOT** need read or write permission for that file.
- ▶ The tests performed by the kernel for file access privilege is a **short-cut** test.



Ownership of New Files and Directories

- ▶ The user ID of a new file is set to the effective user ID of the process.
- ▶ POSIX.1 allows an implementation to choose one of the following options to determine the group ID of a new file:
 - ▶ the effective group ID of the process
 - ▶ the group ID of the directory in which the new file is being created.



access and faccessat Functions I

- ▶ When a process want to test accessibility based on the real user ID and real group ID, `access` function may be useful.

```
1 #include <unistd.h>
2 int access (const char *pathname, int mode);
3 int faccessat(int fd, const char *pathname,
4               int mode, int flag);
```

- ▶ The *mode* is the bitwise OR of any of the constant: `R_OK`, `W_OK`, `X_OK`, `F_OK`.



access and faccessat Functions II

- ▶ The `faccessat` function behaves like `access` when the `pathname` argument is absolute or when the `fd` has the value `AT_FDCWD` and the `pathname` is relative. Otherwise, `faccessat` evaluates the `pathname` relative to the open directory referenced by the `fd`.
- ▶ The `flag` argument can be used to change the behavior of `faccessat`. If the `AT_EACCESS` flag is set, the access checks are made using the effective user and group IDs of the calling process instead of the real user and group IDs
- ▶ Example (Figure 4.8, `filedir/access.c`).



umask Function

- ▶ The **umask** function sets the file mode creation mask for the process and returns the previous value.

```
1 #include <sys/stat.h>
2 mode_t umask(mode_t cmask);
```

- ▶ The **cmask** parameter is formed as the bitwise OR of any of the 9 constants from Figure 4.6.
- ▶ Any bits that are on in the file mode creation mask are turned off in the file's mode.
- ▶ Example (Figure 4.9, **filedir/umask.c**).



chmod, fchmod, and fchmodat Functions I

- ▶ Following functions allow us to change the file access permission for an existing file.

```
1 #include <sys/stat.h>
2 int chmod(const char *pathname, mode_t mode);
3 int fchmod(int fd, mode_t mode);
4 int fchmodat(int fd, const char *pathname,
5              mode_t mode, int flag);
```

- ▶ The *flag* argument can be used to change the behavior of `fchmodat` —when the `AT_SYMLINK_NOFOLLOW` flag is set, `fchmodat` doesn't follow symbolic links.



chmod, fchmod, and fchmodat Functions II

- ▶ To change the permission bits of a file, the effective user ID of the process must equal the owner of the file, or the process must have superuser permission.
- ▶ The *mode* is specified as the bitwise OR of the constants shown in Figure 4.11.
- ▶ Example (Figure 4.12, `filedir/changemod.c`).



Sticky Bit

- ▶ If the sticky bit of a program was set, then the first time the program was executed a copy of the program's text was saved in the swap area when the process terminated. This caused the program to load into memory faster the next time it was executed.
- ▶ If the sticky bit was set for a directory, a file in the directory can be removed or renamed only if the user has write permission for the directory, and either:
 - ▶ owns the file;
 - ▶ owns the directory, or
 - ▶ is the superuser.
- ▶ The sticky bit is **NOT** defined by POSIX.1.



chown, fchown, fchownat, and lchown Functions I

- ▶ These functions allow us to change the user ID and the group ID of a file.

```
1 #include <sys/types.h>
2 #include <unistd.h>
3 int chown (const char *pathname, uid_t owner, gid_t group);
4 int fchown (int fd, uid_t owner, gid_t group);
5 int fchownat(int fd, const char *pathname, uid_t owner,
6               gid_t group, int flag);
7 int lchown (const char *pathname, uid_t owner, gid_t group);
```

- ▶ POSIX.1 allows either form of operation, depending on the value of `_POSIX_CHOWN_RESTRICTED`:
 - ▶ Only the superuser can change the ownership of a file, or
 - ▶ Any user can change the ownership of any files they own.



chown, fchown, fchownat, and lchown Functions II

- ▶ Historically, BSD-based systems have enforced the restriction that only the superuser can change the ownership of a file. This is to prevent users from giving away their files to others, thereby defeating any disk space quota restrictions.
- ▶ If these functions are called by a process other than a superuser process, on successful return, both the set-user-ID and the set-group-ID bits are cleared.



File Size

- ▶ The `st_size` member of the `stat` structure contains the size of the file in bytes. This field is meaningful only for regular files, directories, and symbolic links.
- ▶ `st_blksize` and `st_blocks` means the preferred block size for I/O for the file and the actual number of 512-byte blocks that are allocated.
- ▶ Examples on textbook.



File Truncation

- ▶ Following functions truncate an existing file to *length* bytes, the data beyond *length* is no longer accessible.

```
1 #include <sys/types.h>
2 #include <unistd.h>
3 int truncate (const char *pathname, off_t length);
4 int ftruncate (int fd, off_t length);
```



Filesystems I

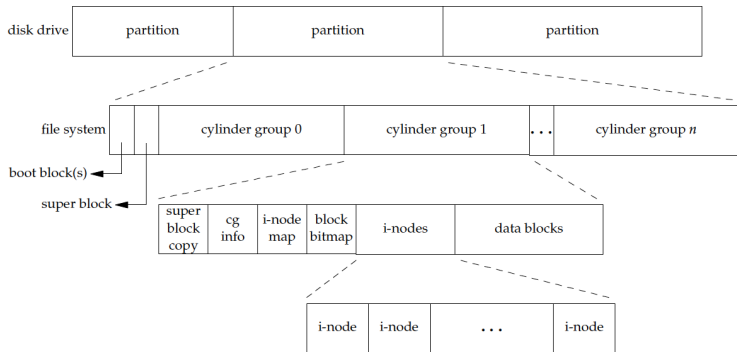


Figure: Disk drive, partitions, and a file system



Filesystems II

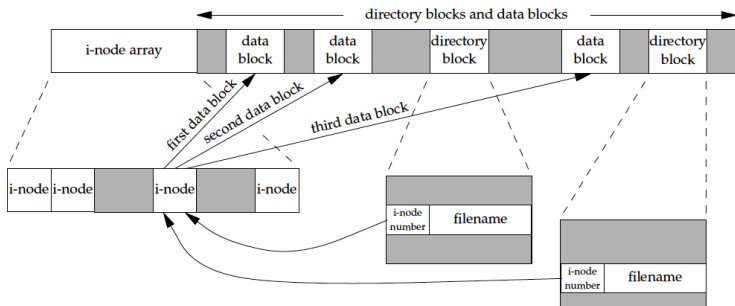


Figure: Cylinder group's i-nodes and data blocks in more detail

Filesystems III

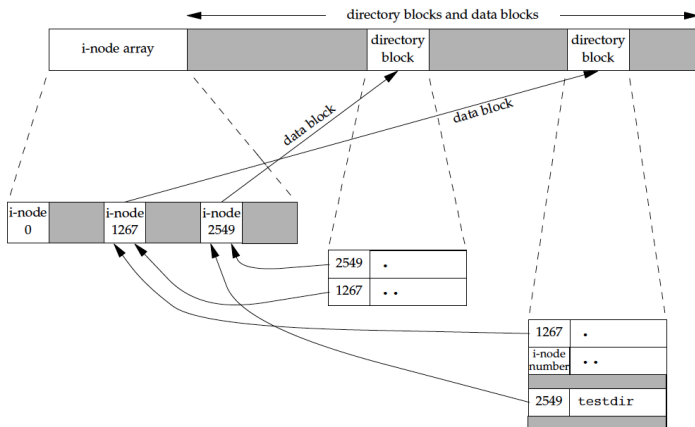


Figure: Sample cylinder group after creating the directory `testdir`

Filesystems IV

- ▶ Every i-node has a link count that contains the number of directory entries that point to the i-node. Only when the link count decrease to 0 can the file be deleted.
- ▶ The `st_nlink` member of `stat` structure contains the link count. The files share the same i-node called **hard link**.
- ▶ The **symbolic link** file has its own i-node. And the actual content of a symbolic file is the pathname of the file that it points to.
- ▶ The i-node contains all the information about the file: the file type, the file's access permission bits, the size of the file, pointers to the data blocks for the file, and so on. The data type for the i-node number is `ino_t`.



Filesystems V

- ▶ A directory entry contains only two items: filename and the i-node number.
- ▶ The actual operation of rename a file is unlink old directory entry and create a new directory entry points to the existing i-node.



link, linkat, unlink, unlinkat, and remove Functions I

- ▶ The `link` and the `linkat` function can create a link to an existing file.

```
1 #include <unistd.h>
2 int link(const char *src, const char *dst);
3 int linkat(int sfd, const char *src,
4            int dfd, const char *dst,
5            int flag);
```

- ▶ The creation of the new directory entry and the increment of the link count must be an `atomic` operation.



link, linkat, unlink, unlinkat, and remove Functions II

- ▶ To remove an existing directory entry we call the `unlink` function:

```
1 #include <unistd.h>
2 int unlink(const char *pathname);
3 int unlinkat(int fd, const char *pathname,
4              int flag);
```

- ▶ If the link count is **NOT** 0 after `unlink` call, the actual data of the file will **NOT** be deleted.
- ▶ As long as some other process has this file open, its contents will not be deleted.
- ▶ Example (Figure 4.16, `filedir/unlink.c`).



link, linkat, unlink, unlinkat, and remove Functions III

- ▶ We can also unlink a file or directory with the `remove` function. For a file, it is identical to `unlink`, for a directory it is identical to `rmdir`.

```
1 #include <stdio.h>
2 int remove(const char *pathname);
```



rename and renameat Functions I

- ▶ A file or a directory is renamed with either the `rename` or `renameat` function.

```
1 #include <stdio.h>
2 int rename(const char *oldname, const char *newname);
3 int renameat(int oldfd, const char *oldname, int newfd,
4             const char *newname);
```

- ▶ Both return: 0 if OK, -1 on error.
- ▶ There are several conditions to describe for these functions, depending on whether *oldname* refers to a file, a directory, or a symbolic link. We must also describe what happens if *newname* already exists.



rename and renameat Functions II

1. If *oldname* specifies a file that is not a directory, then we are renaming a file or a symbolic link. In this case, if *newname* exists, it cannot refer to a directory. Then it been removed, and *oldname* is renamed to *newname*. We must have write permission for the directories containing *oldname* and *newname*.
2. If *oldname* specifies a directory, then we are renaming a directory. If *newname* exists, it must refer to a directory, and that directory must be empty. Then it been removed, and *oldname* is renamed to *newname*. Additionally, when we're renaming a directory, *newname* cannot contain a path prefix that names *oldname*.
3. If either *oldname* or *newname* refers to a symbolic link, then the link itself is processed, not the file to which it points.



rename and renameat Functions III

4. We can't rename dot or dot-dot. More precisely, neither dot nor dot-dot can appear as the last component of *oldname* or *newname*.
 5. As a special case, if *oldname* and *newname* refer to the same file, the function returns successfully without changing anything.
- If *newname* already exists, we need permissions as if we were deleting it. Also, because we're removing the directory entry for *oldname* and possibly creating a directory entry for *newname*, we need write permission and execute permission in the directories containing *oldname* and *newname*.



rename and renameat Functions IV

- ▶ The `renameat` function provides the same functionality as the `rename` function, except when either *oldname* or *newname* refers to a relative pathname. In this case, it is evaluated relative to the directory referenced by *oldfd*. Similarly, *newname* is evaluated relative to the directory referenced by *newfd* if *newname* specifies a relative pathname. Either the *oldfd* or *newfd* arguments (or both) can be set to `AT_FDCWD` to evaluate the corresponding pathname relative to the current working directory.



Symbolic Link

- ▶ **Symbolic links** were introduced to get around the limitations of hard links:
 - ▶ Hard links normally require that the link and the file reside in the same filesystem;
 - ▶ Only the superuser can create a hard link to a directory.
- ▶ Symbolic link are typically used to move a file or an entire directory hierarchy to some other location on a system.
- ▶ Figure 4.17 summarizes whether the function described in this chapter follow a symbolic link or not.
- ▶ Whether **chown** follows a symbolic link or not depends on the implementation.
- ▶ Example: a loop of symbolic link (Figure 4.18).



Creating and Reading Symbolic Links

- ▶ A symbolic link is created with the `symlink` or `symlinkat` function.

```
1 #include <unistd.h>
2 int symlink(const char *src, const char *dst);
3 int symlinkat(const char *src, int fd, const char *dst);
```

- ▶ Since the `open` function follows a symbolic link, we need a way to open the link itself and read the name in the link. The `readlink` and `readlinkat` functions do this.

```
1 #include <unistd.h>
2 ssize_t readlink(const char* restrict pathname,
3                 char* restrict buf, size_t bufsz);
4 ssize_t readlinkat(int fd, const char * restrict pathname,
5                   char* restrict buf, size_t bufsz);
```

- ▶ This functions combines the actions of `open`, `read`, and `close`.



File Time

- ▶ Three time fields are maintained for each file:
 - `ST_ATIME`: last access time of file data;
 - `ST_MTIME`: last modification time of file data;
 - `ST_CTIME`: last change time of i-node status.
- ▶ The system does not maintain the last access time for an **i-node**.
- ▶ The `ls(1)` command displays or sorts only on one of the three time value. By default **LMT**, the `-u` option means **LAT**, and the `-c` means **LCT**.
- ▶ Figure 4.20 summarizes the effects of the various functions that we've described on these three times.



futimens, utimesat, and utime Functions I

- ▶ The LAT and the LMT of a file will be set to the greatest value supported by the filesystem that is not greater than the specified time with these functions in nanosecond precision:

```
1 #include <sys/stat.h>
2 int futimens(int fd, const struct timespec times[2]);
3 int utimensat(int dirfd, const char *path,
4               const struct timespec times[2],
5               int flags);
6
7 struct timespec {
8     time_t tv_sec;    /* seconds */
9     long   tv_nsec;   /* nanoseconds */
10 };
```

- ▶ In both functions, the 1st element of the *times* array contains the LAT, and the 2nd element contains the LMT. The two time values are calendar times.



futimens, utimesat, and utime Functions II

- ▶ Timestamps can be specified in one of four ways:
 1. The *times* argument is a null pointer. In this case, both timestamps are set to the current time.
 2. If either *tv_nsec* field has the special value *UTIME_NOW*, the corresponding timestamp is set to the current time. The corresponding *tv_sec* field is ignored.
 3. If either *tv_nsec* field has the special value *UTIME_OMIT*, then the corresponding timestamp is left unchanged. The corresponding *tv_sec* field is ignored.
 4. The *times* argument points to an array of two timespec structures and the *tv_nsec* field contains a value other than *UTIME_NOW* or *UTIME_OMIT*. Then the corresponding timestamp is set to the value specified by the corresponding *tv_sec* and *tv_nsec* fields.



futimens, utimesat, and utime Functions III

- ▶ The privileges required to execute these functions depend on the value of the *times* argument.
- ▶ Both *futimens* and *utimensat* are included in POSIX.1. A third function, *utimes*, is included in the Single UNIX Specification as part of the XSI option.

```
1 #include <sys/time.h>
2 int utimes (const char *pathname,
3             const struct timeval times[2]);
4
5 struct timeval {
6     time_t tv_sec;    /* seconds */
7     long   tv_usec;   /* microseconds */
8 };
```



futimens, utimesat, and utime Functions IV

- ▶ The `utimes` function operates on a pathname. The *times* argument is expressed in seconds and microseconds.
- ▶ Note that we are unable to specify a value for the changed-status time as this field is automatically updated when the `utime` functions family are called.
- ▶ Example (Figure 4.21, `filedir/zap.c`).



mkdir, mkdirat, and rmdir Functions

- ▶ Directories are created and deleted with `mkdir` and `rmdir` functions.

```
1 #include <sys/stat.h>
2 int mkdir(const char *pathname, mode_t mode);
3 int mkdirat(int fd, const char *pathname,
4             mode_t mode);
```

- ▶ Be careful of *mode* parameter.

```
1 #include <unistd.h>
2 int rmdir (const char *pathname);
```

- ▶ The *pathname* parameter must be an empty directory.
- ▶ When this function is called, the *pathname* will be locked.



Reading Directories I

- ▶ Directories can be read by anyone who has access permission to read the directory. But only kernel can write to a directory.
- ▶ A set of directory routines were developed and are part of POSIX.1.

```
1 struct dirent{
2     ino_t    d_ino;
3     char     d_name[256];
4 };
5
6 #include <dirent.h>
7 DIR *opendir(const char *pathname);
8 DIR *fdopendir(int fd);
9 struct dirent *readdir(DIR *dp);
```



Reading Directories II

```
10 void rewinddir(DIR *dp);  
11 int closedir(DIR *dp);  
12 long telldir(DIR *dp);  
13 void seekdir(DIR *dp, long loc);
```

- ▶ The `telldir` and `seekdir` functions are not part of the base POSIX.1 standard. They are included in the XSI option in the SUS, so all conforming UNIX System implementations are expected to provide them.
- ▶ The ordering of entries within the directory is implementation dependent and is usually not alphabetical.
- ▶ Example (Figure 4.22, `filedir/ftw8.c`).



chdir, fchdir, and getcwd Functions I

- ▶ Every process has a current working directory.
- ▶ We can change the current working directory of the calling process by calling the `chdir` or `fchdir` functions.

```
1 #include <unistd.h>
2 int chdir (const char *pathname);
3 int fchdir (int fd);
```

- ▶ Both return 0 if OK, -1 on error.
- ▶ Since the current working directory is an attribute of a process it cannot affect others processes.
- ▶ Example (Figure 4.23, `filedir/mycd.c`).



chdir, fchdir, and getcwd Functions II

- ▶ Unfortunately, all the kernel maintains for each process is the i-node number and device id for the CWD. So we need a function that can obtain the entire absolute pathname of the CWD by the giving i-node number.

```
1 #include <unistd.h>  
2 char *getcwd (char *buf, size_t size);
```

- ▶ Example (Figure 4.24, `filedir/cdpwd.c`).



Device Special Files I

- ▶ Every filesystem is known by its major and minor device number. This device number is encoded in the primitive system data type `dev_t`.
- ▶ The major number identifies the device driver and sometimes encodes which peripheral board to communicate with; the minor number identifies the specific subdevice.
- ▶ We can usually access the major and minor device numbers through two macros: `major` and `minor`.
- ▶ Linux store the device number in a 16-bit integer with higher 8 bits for the major and lower 8 bits for minor device number.
- ▶ These two macros are defined in `sys/sysmacros.h`.



Device Special Files II

- ▶ The `st_dev` value for every filename on a system is the device number of the filesystem containing that filename and its i-node.
- ▶ Only character and block special device files have an `st_rdev` value. This value contains the device number for the actual device.
- ▶ Example (Figure 4.25, `filedir/devrdev.c`).



Summary

- ▶ This chapter centered around the `stat` function. We've gone through each member in the `stat` structure in detail.
- ▶ A through understanding of all the properties of files and directories and all the functions that operate on them is essential to UNIX programming.



The End

The End of Chapter 4.

