

# Advanced Programming in the UNIX Environment — *Process Environment*

Hop Lee  
hoplee@bupt.edu.cn

SCHOOL OF INFORMATION AND COMMUNICATION ENGINEERING



# Table of Contents

MAIN Function

Process Termination

Exit Functions

atexit Function

Command-Line Arguments

Environment List

Memory Layout of a C Program

Shared Library

Memory Allocation

Environment Variables

SETJMP and LONGJMP Functions

GETRLIMIT and SETRLIMIT Functions



## main Function

- ▶ A C program starts execution with a function called `main`:

```
1 int main(int argc, char *argv[]);
```

- ▶ `argc` is the number of command-line arguments and `argv` is an array of pointers to the arguments.
- ▶ When a C program is started by the kernel (by one of the `exec` functions), a special **start-up routine** is called before the `main` function is called.
- ▶ The executable program file specifies this start-up routine as the starting address for the program. This start-up routine takes values from the kernel and sets things up so that the `main` function is called as shown earlier.



# Process Termination

- ▶ There are eight ways for a process to terminate:
  1. Normal termination
    - 1.1 return from `main`
    - 1.2 calling `exit`
    - 1.3 calling `_exit` or `_Exit`
    - 1.4 Return of the last thread from its start routine (Section 11.5)
    - 1.5 Calling `pthread_exit` (Section 11.5) from the last thread
  2. Abnormal termination
    - 2.1 calling `abort`
    - 2.2 terminated by a signal
    - 2.3 Response of the last thread to a cancellation request (Sections 11.5 and 12.7)
- ▶ The start-up routine that we mentioned in the previous section is also written so that if the `main` function returns, the `exit` function is called:

```
1 exit(main(argc, argv));
```



# Exit Functions I

- ▶ These functions terminate a program normally:
  - ▶ `_exit`, which returns to the kernel immediately
  - ▶ `exit`, which performs certain cleanup processing and then returns to the kernel.

```
1 #include <stdlib.h>
2 void exit(int status);
3 void _Exit(int status);
4 #include <unistd.h>
5 void _exit(int status);
```

- ▶ Both these functions expect a single integer argument, which we call the **exit status**.
- ▶ If



## Exit Functions II

1. either of these functions is called without an exit status
2. `main` does a `return` without a return value
3. `main` falls off the end (an implicit return)

the exit status of the process is undefined.

- ▶ Returning an integer value from the `main` function is equivalent to calling `exit` with the same value.
- ▶ Example (Figure 7.1, `environ/hello1.c`).



## atexit Function I

- ▶ With ISO C a process can register up to 32 functions that are automatically called by `exit`. These so called **exit handlers** are registered by `atexit` function.

```
1 #include <stdio.h>  
2 int atexit(void (*func)(void));
```

- ▶ When exit handler is called it is not passed any arguments and it is not expected to return a value.
- ▶ The `exit` function calls these functions in reverse order of their registration. Each function is called as many times as it was registered.
- ▶ Example (Figure 7.3, `environ/doatexit.c`).



## atexit Function II

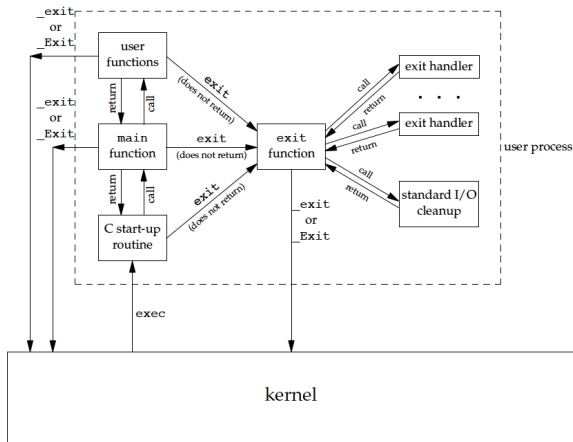


Figure: How a C program is started and how it terminates



# Command-Line Arguments

- ▶ When a program is executed, the process that does the `exec` can pass command-line arguments to the new program.
- ▶ Example (Figure 7.4, `environ/echoarg.c`).
- ▶ We are guaranteed by both ISO C and POSIX.1 that `argv[argc]` is a null pointer. This let us alternatively code the argument processing loop as

```
1 for(i = 0; argv[i] != NULL; i++)
```



## Environment List I

- ▶ Each program is also passed an **environment list**.
- ▶ The environment list is an array of character pointers, with each pointer containing the address of a null-terminated C string. The address of the array of pointers is contained in the global variable `environ`

```
1 extern char **environ
```



## Environment List II

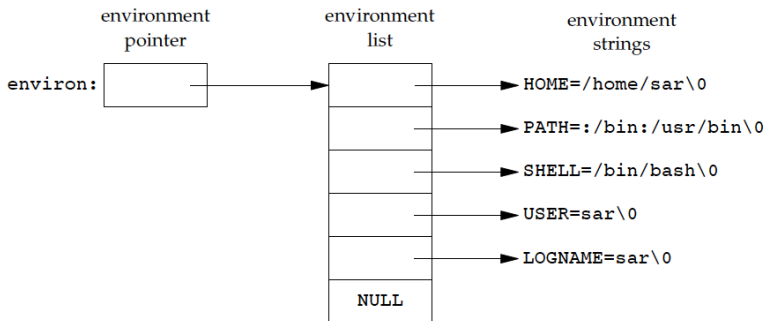


Figure: Environment consisting of five C character strings

# Memory Layout of a C Program I

- ▶ Historically a C program has been composed of the following pieces:

**Text segment** The machine instructions that are executed by the CPU. Usually the text segment is sharable and read only.

**Initialized data segment** Also called **data segment** and it contains variables that are specifically initialized in the program.

**Uninitialized data segment** Also called **bss segment**, named after “block started by symbol”. Data in this segment is initialized by the kernel to arithmetic 0 or null pointers before the program starts executing.



# Memory Layout of a C Program II

**Stack** This is where automatic variables are stored, along with information that is saved each time a function is called.

**Heap** Dynamic memory allocation usually takes place on the heap. Historically the heap has been located between the top of the uninitialized data and the bottom of the stack.

**Arguments and Environments** This area contains command-line arguments and working environment-list.

- ▶ The **size**(1) command reports the sizes (in bytes) of the text, data, and bss segments.



## Memory Layout of a C Program III

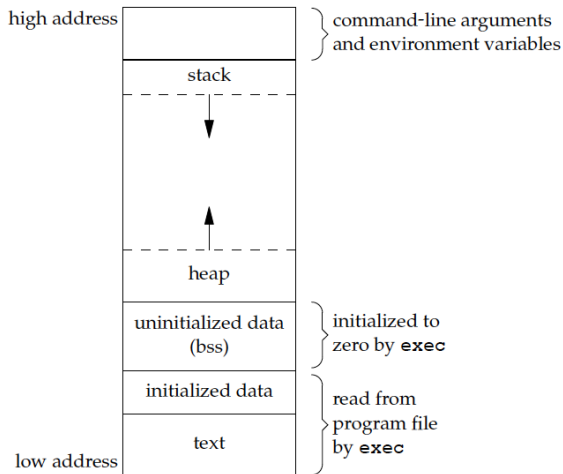


Figure: Typical memory arrangement



## Shared Library

- ▶ Most modern UNIX implementation support dynamic shared library.
- ▶ This technology make it possible that the executable file need NOT include the popular library functions, but put a copy of these routines in an area which is readable for all process.
- ▶ This method reduce the length of program file, but increase the run time cost.
- ▶ Another advantage of shared libraries is that library functions can be replaced with new versions without having to relink edit every program that uses the library
- ▶ Under Linux, `gcc`, use `-static` option will force the program to be static linked.



# Memory Allocation I

- ▶ There are three functions specified by ISO C for memory allocation.
  - ▶ `malloc`. Allocates a specified number of bytes of memory. The initial value of the memory is indeterminate.
  - ▶ `calloc`. Allocates space for a specified number of objects of a specified size. The space is initialized to all bit 0.
  - ▶ `realloc`. Change the size of a previously allocated area (increases or decreases).

```
1 #include <stdlib.h>
2 void *malloc(size_t size);
3 void *calloc(size_t nobj, size_t size);
4 void *realloc(void *ptr, size_t newsize);
5 void free(void *ptr);
```





## Memory Allocation II

- ▶ The function `free` causes the space pointed to by `ptr` to be de-allocated.
- ▶ Many replacements for `malloc` and `free` are available. Some systems already include libraries providing alternative memory allocator implementations. Other systems provide only the standard allocator, leaving it up to software developers to download alternatives, if desired.



# Environment Variables I

- ▶ All the environment variable looks like this:

```
1 name=value
```

- ▶ The UNIX kernel do not care the meaning of the value. It's explanation depends on the application program.
- ▶ ISO C define the `getenv` function:

```
1 #include <stdlib.h>  
2 char *getenv(const char *name);
```

It return a pointer to the *value* string.

- ▶ We should always access the environment variable through the `getenv` function but not the global variable `environ`.



## Environment Variables II

- ▶ There are also some other functions used to create new environment variable or alert the value of an exist environment variable or delete certain environment variable.

```
1 #include <stdlib.h>
2 int putenv(char *string);
3 int setenv(const char *name, const char *value,
4           int overwrite);
5 void unsetenv(const char *name);
6 int clearenv(void);
```



## Environment Variables III

- ▶ The `putenv` function adds or changes the value of environment variables. The argument *string* is of the form *name=value*. If *name* does not already exist in the environment, then string is added to the environment. If *name* does exist, then the value of *name* in the environment is changed to *value*. The string pointed to by *string* becomes part of the environment, so altering the string changes the environment.
- ▶ The `setenv` function adds the variable name to the environment with the value *value*, if *name* does not already exist. If *name* does exist in the environment, then its value is changed to *value* if *overwrite* is non-zero; if *overwrite* is zero, then the value of *name* is not changed.



## Environment Variables IV

- ▶ The `unsetenv` function deletes the variable *name* from the environment.
- ▶ The `clearenv` function clears the environment of all *name-value* pairs and sets the value of the external variable `environ` to `NULL`.



## setjmp and longjmp Functions I

- ▶ We can not use `goto` statement across the function in C language. But we can use nonlocal goto functions `setjmp` and `longjmp` instead.
- ▶ They are useful for dealing with errors and interrupts encountered in a deeply nested function call of a program.
- ▶ Example (Figure 7.9, `environ/cmd1.c`).

```
1 #include <setjmp.h>
2 int setjmp(jmp_buf env);
3 void longjmp(jmp_buf env, int val);
```

- ▶ `longjmp` restores the environment saved by the last call of `setjmp` with the corresponding `env` argument.



## setjmp and longjmp Functions II

- ▶ After `longjmp` is completed, program execution continues as if the corresponding call of `setjmp` had just return the value `val`.
- ▶ Example (Figure 7.11, `environ/cmd2.c`).
- ▶ Example (Figure 7.13, `environ/testjmp.c`).
- ▶ Potential problem with automatic variables. The basic rule is that an automatic variable can never be referenced after the function that declared it returns. Figure 7.14 (`environ/opendata.c`) shows an example.



## getrlimit and setrlimit Functions I

- ▶ Every process has a group of resource limits, some of them can be queried and modified by `getrlimit` and `setrlimit` functions.

```
1 #include <sys/resource.h>
2 int getrlimit(int resource, struct rlimit *rlim);
3 int setrlimit(int resource,
4               const struct rlimit *rlim);
```

- ▶ Each resource has an associated soft and hard limit, as defined by the `rlimit` structure:





## getrlimit and setrlimit Functions II

```
1 struct rlimit {  
2     rlim_t rlim_cur; /* Soft limit */  
3     rlim_t rlim_max; /* Hard limit (ceiling  
4                     for rlim_cur) */  
5 };
```

- ▶ Three rules govern the changing of the resource limits.
  1. A process can change its soft limit to a value less than or equal to its hard limit.
  2. A process can lower its hard limit to a value greater than or equal to its soft limit. This lowering of the hard limit is irreversible for normal users.
  3. Only a superuser process can raise a hard limit.
- ▶ An infinite limit is specified by the constant `RLIM_INFINITY`.



## getrlimit and setrlimit Functions III

- ▶ The *resource* indicate one of the pre-defined resource name shown in Figure 7.15.
- ▶ The resource limits affect the calling process and are inherited by any of its children.
- ▶ Example (Figure 7.16, `environ/getrlimit.c`).
- ▶ Note that we've used the ISO C string-creation operator (`#`) in the `doit` macro, to generate the string value for each resource name.



## The End

The End of Chapter 7.

