

Advanced Programming in the UNIX Environment — *Process Control*

Hop Lee
hoplee@bupt.edu.cn

Contents

1	Process Identifiers	1
2	fork Function	2
3	vfork Function	4
4	exit Function	4
5	wait and waitpid Functions	5
6	waitid Function	6
7	wait3 and wait4 Functions	6
8	Race Conditions	6
9	exec Function Family	7
10	Changing User IDs and Group IDs	8
10.1	setreuid and setregid Functions	9
10.2	seteuid and setegid Functions	9
11	Interpreter Files	9
12	system Function	10
12.1	Set-User-ID Programs	10
13	Process Accounting	11
14	User Identification	12
15	Process Scheduling	12
16	Process Times	13

1 Process Identifiers

- Every UNIX process is guaranteed to have a unique non-negative integer as numeric identifier called **process ID**.

- Example (Figure 1.6, `intro/hello.c`).
- Process ID 0 is usually scheduler process and is often known as the `swapper`. No program on disk corresponds to this process, which is part of the kernel and is known as a system process.
- Process ID 1 is usually the `init`¹ process and is invoked by the kernel at the end of the bootstrap procedure. The program file for this process is `/sbin/init`. The `init` process never dies. It is a normal user process, not a system process within the kernel, although it does run with superuser privileges.
- Except process ID, there are several other identifications for each process. We can get these identifications by such functions:

```
1 #include <unistd.h>
2 pid_t getpid(void);
3 pid_t getppid(void);
4 uid_t getuid(void);
5 uid_t geteuid(void);
6 gid_t getgid(void);
7 gid_t getegid(void);
```

- These functions are always successful.

2 fork Function

- An existing process can create a new one by calling the `fork` function.

```
1 #include <unistd.h>
2 pid_t fork(void);
```

- The new process created by `fork` is called the **child process**.
- This function is called once but returns twice. The only difference in the returns is that the return value in the child is 0, whereas in the parent is the process ID of the new child.
- Both the child and the parent continue executing with the instruction that follows the call to `fork`. The child is a copy of the parent. Note that this is a copy for the child; the parent and the child do not share these portions of memory. The parent and the child do share the text segment, however.
- In modern implementations, `fork` is implemented using copy-on-write technique, so the only penalty incurred by `fork` is the time and memory required to duplicate the parent's page tables, and to create a unique task structure for the child.
- Variations of the `fork` function are provided by some platforms.
- Example (Figure 8.1, `proc/fork1.c`).
- All file descriptors that are opened in the parent are duplicated in the child after `fork`.
- Besides the open files, numerous other properties of the parent are inherited by the child:
 - ruid, rgid, euid and egid
 - Supplementary group IDs

¹ `launchd` under MacOS 10.10.2

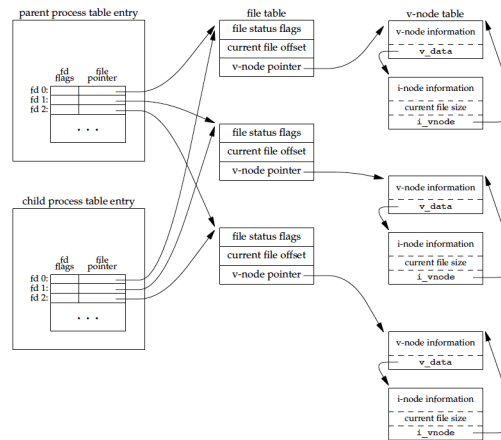


Figure 1: Sharing of open files between parent and child after `fork`

- Process group ID
 - Session ID
 - Controlling terminal
 - Current working directory
 - Root directory
 - set-user-ID and set-group-ID flags
 - File mode creation mask
 - Signal mask and dispositions
 - The close-on-exec flag for any open file descriptors
 - Environment
 - Attached shared memory segments
 - Memory mappings
 - Resource limits
- The differences between the parent and child are
 - The return values from `fork`.
 - The pids.
 - The ppids.
 - The child's `tms_utime`, `tms_stime`, `tms_cutime`, and `tms_cstime` values are set to 0 (these times are discussed in Section 8.17).
 - File locks set by the parent are not inherited by the child.
 - Pending alarms are cleared for the child.
 - The set of pending signals for the child is set to the empty set.
 - The two main reasons for `fork` fail are
 1. if too many processes there are already in the system
 2. if the total number of processes for this real user ID exceeds the system's limit.

- `fork` has two main usage:
 1. When a process wants to duplicate itself so that the parent and child can each execute different section of code at the same time. This is common for network servers.
 2. When a process wants to execute a different program. This is common for shells. In this case the child does an `exec` calling right after it returns from the `fork`.

3 vfork Function

- The `vfork` function has the same calling sequence and same return values as `fork`.

```
1 #include <unistd.h>
2 pid_t vfork(void);
```

- `vfork` creates the new process without copying the address space of the parent into the child. And the child just
 1. calling the `exec` or `exit` right after the `vfork`
 2. running in the address space of the parent until it calls either `exec` or `exit`
- This optimization is more efficient on some implementations of the UNIX System, but leads to undefined results if the child modifies any data, makes function calls, or returns without calling `exec` or `exit`.
- Another different between the two functions is that `vfork` guarantees that the child runs first. This feature can lead to *deadlock* if the child depends on further actions of the parent.
- The `vfork` and the `fork` system calls are identity in Linux.
- Example (Figure 8.3, `proc/vfork1.c`).
- Note in Figure 8.3 that we call `_exit` instead of `exit`.

4 exit Function

- Regardless how a process terminates, the same code in the kernel is eventually executed. This kernel code closes all the open descriptors for the process, release memory that it was using, and so on.
- We want the parent to be notifying how its child terminated. For the three exit functions this is done by passing an **exit status** as the argument to the function. For an abnormal termination, however, the kernel generates a **termination status** to indicate the reason for the abnormal termination.
- If the parent terminates before the child, the `init` process will becomes the parent process of the orphan. Then the terminal status of the child will be collected by the `init` process.
- If the child terminates before the parent, the kernel will keep a certain amount of information for every terminating process, so that the information is available when the parent calls `wait` or `waitpid`. After this operation, the kernel can discard all the memory used by the process and close its open files.
- In UNIX terminology the process that has terminated, but whose parent has NOT yet waited for it, is called a **zombie**.
- When a process who's parent is `init` terminates, it will definitely *NOT* become a zombie.

5 wait and waitpid Functions

- A process that calls `wait` or `waitpid` can
 - block (if all of its children are still running), or
 - return immediately with the termination status of a child (if a child has terminated and is waiting for its termination status to be fetched), or
 - return immediately with an error (if it does *NOT* have any child processes)
- If the process calling `wait` because it received the `SIGCHLD` signal, we expect `wait` to return immediately. But if we call it at any random point in time, it can block.

```
1 #include <sys/types.h>
2 #include <sys/wait.h>
3 pid_t wait(int *statloc);
4 pid_t waitpid(pid_t pid, int *statloc,
5               int options);
```

- Both will return the child's pid when success.
- The differences between these two functions are
 - `wait` can block the caller until a child process terminates, while `waitpid` has an option that prevents it from blocking,
 - `waitpid` doesn't wait for the first child to terminate; it has a number of options that control which process it waits for.
- For both functions the argument `statloc` is a pointer to an integer which will store the termination status of the terminated child process. If we do not care the termination status, we just pass a null pointer as this argument.
- There are four mutually exclusive macros that tell us how the process terminated. Based on which of these three macros is true, other macros are used to obtain the exit status, signal number, and the like.
- Example (Figure 8.5, `lib/prexit.c`).
- The interpretation of the `pid` argument for `waitpid` depends on its value:

<code>pid == -1</code>	waits for any child process
<code>pid > 0</code>	waits for the child whose process ID equals <code>pid</code>
<code>pid == 0</code>	waits for any child whose process group ID equals that of the calling process
<code>pid < -1</code>	waits for any child whose process group ID equals the absolute value of <code>pid</code>
- Example (Figure 8.6, `proc/wait1.c`).
- The `options` argument lets us further control the operation of `waitpid`. This argument either is 0 or is constructed from the bitwise OR of the constants in Figure 8.7.
- The `waitpid` function provides three features that aren't provided by the `wait` function.
 1. The `waitpid` function lets us wait for one particular process, whereas the `wait` function returns the status of any terminated child.
 2. The `waitpid` function provides a nonblocking version of `wait`.
 3. The `waitpid` function provides support for job control with the `WUNTRACED` and `WCONTINUED` options.
- Example (Figure 8.8, `proc/fork2.c`).

6 waitid Function

- The SUS includes an additional function to retrieve the exit status of a process. The `waitid` function is similar to `waitpid`, but provides extra flexibility.

```
1 #include <sys/wait.h>
2 int waitid(idtype_t idtype, id_t id,
3           siginfo_t *infop, int options);
```

- `waitid` allows a process to specify which children to wait for. Instead of encoding this information in a single argument combined with the process ID or process group ID, two separate arguments are used. The `id` parameter is interpreted based on the value of `idtype`. The types supported (`P_PID`, `P_PGID`, and `P_ALL`) are summarized in Figure 8.9.
- The `options` argument is a bitwise OR of the flags shown in Figure 8.10. These flags indicate which state changes the caller is interested in.
- At least one of `WCONTINUED`, `WEXITED`, or `WSTOPPED` must be specified in the `options` argument.
- The `infop` argument is a pointer to a `siginfo` structure. This structure contains detailed information about the signal generated that caused the state change in the child process.

7 wait3 and wait4 Functions

- Most UNIX system implementations provide two additional functions, `wait3` and `wait4`. These two functions allow the kernel to return a summary of the resources used by the terminated process and all its child processes.

```
1 #include <sys/types.h>
2 #include <sys/wait.h>
3 #include <sys/time.h>
4 #include <sys/resource.h>
5 pid_t wait3(int *statloc, int options,
6            struct rusage *rusage);
7 pid_t wait4(pid_t pid, int *statloc, int options,
8            struct rusage *rusage);
```

- The resource information includes information such as the amount of user CPU time, amount of system CPU time, number of page faults, number of signals received, and the like. Refer to the `getrusage(2)` manual page for additional details.

8 Race Conditions

- A **race condition** occurs when multiple processes are trying to do something with shared data and the final outcome depends on the order in which the processes run.
- To avoid race conditions and to avoid polling, some form of signaling is required between multiple processes.
- Example (Figure 8.12, `proc/tellwait1.c`), program with a race condition.
- Example (Figure 8.13, `proc/tellwait2.c`), modification of Figure 8.12 to avoid race condition.

9 exec Function Family

- When a process calls one of the `exec` functions, that process is completely replaced by the new program, and the new program starts executing at its `main` function.
- `exec` replaces the current process (its text, data, heap, and stack segments) with a brand new program from disk.

```
1 #include <unistd.h>
2 int execl(const char *pathname, const char *arg0, ...);
3 int execv(const char *pathname, char *const argv[]);
4 int execlp(const char *pathname, const char *arg0, ...,
5           /* (char *)0, char *const envp[] */);
6 int execve(const char *pathname, char *const argv[],
7           char *const envp[]);
8 int execlp(const char *filename, const char *arg0, ...);
9 int execvp(const char *filename, char *const argv[]);
10 int fexecve(int fd, char *const argv[],
11            char *const envp[]);
```

- When a `filename` argument is specified
 - if `filename` contains a slash, it is taken as a pathname
 - otherwise, the executable file is searched for in the directories specified by the `PATH` environment variableand if the found file is not a machine executable that was generated by the `link` utility, it assumes the file is a shell script and tries to invoke `/bin/sh` with the `filename` as input to the shell.
- With `fexecve`, we avoid the issue of finding the correct executable file altogether and rely on the caller to do this. By using a file descriptor, the caller can verify the file is in fact the intended file and execute it without a race.
- The function `execl`, `execlp` and `execvp` require each of the command-line arguments to the new program to be specified as separate arguments with the end as a null pointer. For the other four functions, we have to build an array of pointers to the arguments, and the address of this array is the argument to these three functions.
- The three functions whose name ended in an `e` allow us to pass a pointer to an array of pointers to the environment strings. The other four functions, however, use the `environ` variable in the calling process to copy the existing environment for the new program.
- Example (Figure 8.16, `proc/exec1.c`, `proc/echoall.c`).
- Only one of these seven functions, `execve`, is a system call within the kernel. The other six are just library functions that eventually invoke this system call.
- The process ID does not change after an `exec`, but the new program inherits additional properties from the calling process:
 - pid and ppid
 - ruid and rgid
 - Supplementary gids
 - Process group ID
 - Session ID
 - Root directory

- File locks
- Resource limits
- Current working directory
- Controlling terminal
- Time left until alarm clock
- File mode creation mask
- Process signal mask
- Pending signals
- Nice value
- Values for `tms_utime`, `tms_stime`, `tms_cutime`, and `tms_cstime`

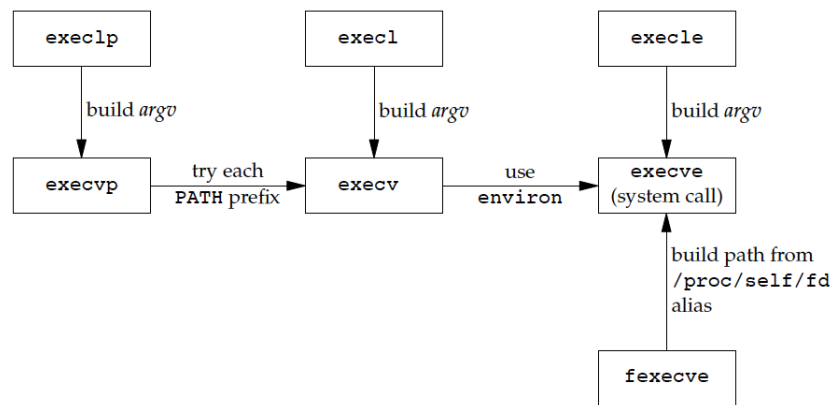


Figure 2: Relationship of the seven exec functions

10 Changing User IDs and Group IDs

- We can set the real user ID and effective user ID with the `setuid` function. Similarly, the `setgid` function.

```

1 #include <sys/types.h>
2 #include <unistd.h>
3 int setuid(uid_t uid);
4 int setgid(gid_t gid);

```

- There are rules for who can change the IDs:
 - If the process has superuser privileges, the `setuid` functions sets the real user ID, effective user ID, and saved set-user-ID to `uid`.
 - If the process does not have superuser privileges, but `uid` equals either the real user ID or the saved set-user-ID of the process, `setuid` sets only the effective user ID to `uid`.
 - If neither of these two conditions is true, `errno` is set to `EPERM` and an error is returned.
- There are several statements about the three user IDs:
 - Only a superuser process can change the real user ID.

- The effective user ID is set by the `exec` functions, only if the set-user-ID bit is set for the program file.
- We can call `setuid` at any time to set the effective user ID to either the real user ID or the saved set-user-ID.
- The saved set-user-ID is copied from the effective user ID by `exec` after the `exec` stores the effective user ID from the file's user ID.
- We can not obtain the current value of the saved set-user-ID.

10.1 setreuid and setregid Functions

- 4.3+BSD supports the swapping of the real user ID and the effective user ID with the following functions.

```
1 #include <sys/types.h>
2 #include <unistd.h>
3 int setreuid(uid_t ruid, uid_t euid);
4 int setregid(gid_t rgid, gid_t egid);
```

10.2 seteuid and setegid Functions

- A proposed change to POSIX.1 includes the two functions which only change the effective user ID or effective group ID:

```
1 #include <sys/types.h>
2 #include <unistd.h>
3 int seteuid(uid_t euid);
4 int setegid(gid_t egid);
```

- Example, p259.

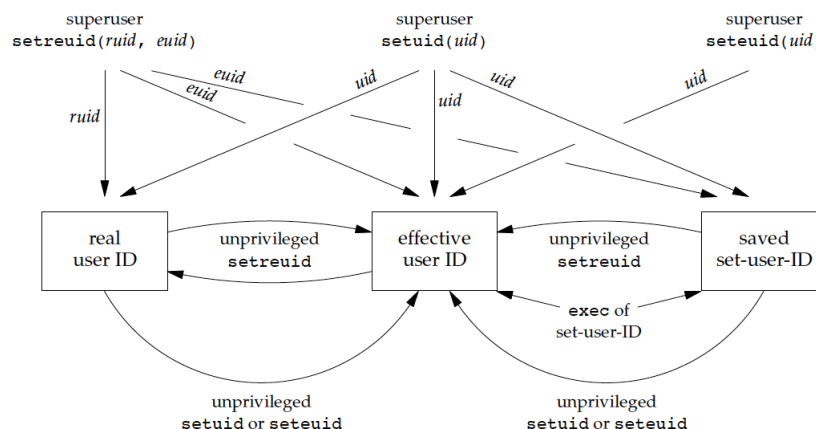


Figure 3: Summary of all the functions that set the various user IDs

11 Interpreter Files

- The another name of 'interpreter files' is **script**.
- Both SVR4 and 4.3+BSD support interpreter files. These files are text files that begin with a line of the form

```
1 #! pathname [optional-argument]
```

- The *pathname* is normally an absolute pathname since the environment variable `PATH` is not used.
- The actual file that gets `exec`ed by the kernel is not the script file, but the file specified by the *pathname* on the first line of the interpreter file, interpreter.
- Example (Figure 8.20, `proc/exec2.c`, `proc/testinterp`, `environ/echoarg.c`).
- The corresponding output is:

```
1 [hop@Hanoi proc]$ cat testinterp
2 #!/home/hop/bin/echoarg foo
3 [hop@Hanoi proc]$ ./exec2
4 argv[0]: /home/hop/bin/echoarg
5 argv[1]: foo
6 argv[2]: /home/hop/bin/testinterp
7 argv[3]: myarg1
8 argv[4]: MY ARG2
```

12 system Function

- ISO C defines the `system` function, but its operation is strongly system dependent.

```
1 #include <stdlib.h>
2 int system(const char *string);
```

- `system` executes a command specified in *string* by calling `/bin/sh -c string`, and returns after the command has been completed. During execution of the command, `SIGCHLD` will be blocked, and `SIGINT` and `SIGQUIT` will be ignored.
- The value returned is -1 (`fork` or `waitpid` fails) or 127 (`exec` fails) on error, and the return status of the command otherwise. This latter return status is in the format specified in `waitpid`. Thus, the exit code of the command will be `WEXITSTATUS(status)`.
- In case `/bin/sh` could not be executed, the exit status will be that of a command that does `exit(127)`.
- If the value of *string* is `NULL`, `system` returns nonzero if the shell is available, and zero if not.
- `system` does not affect the wait status of any other children.
- Figure 8.22 is an implementation of the `system` function. And Figure 8.23 is a test program.
- Example (Figure 8.22, 8.23, `proc/system1.c`, `proc/systest1.c`).
- The advantage in using `system`, instead of using `fork` and `exec` directly, is that `system` does all the required error handling and all the required signal handling.

12.1 Set-User-ID Programs

- Calling `system` from a set-user-ID program is a security hole and should never be done.
- The superuser permissions that we gave to the set-user-ID program are retained across the `fork` and `exec` that are done by `system`.
- Example (Figure 8.24, 8.25, `proc/systest3.c`, `proc/pruids.c`).

13 Process Accounting

- Most UNIX systems provide an option to do process accounting.
- Process accounting is NOT specified by any of the standards.
- There is a command `accton` under SVR4 and 4.3+BSD. A superuser can execute it with a pathname to the log file as the argument.
- The structure of the accounting records is defined in the header `sys/acct.h`.
- The `ac_flag` member records certain events during the execution of the process. There are described in Table 1:

Table 1: Values for `ac_flag` from accounting record.

<code>ac_flag</code>	Description
AFORK	process is the result of <code>fork</code> , but never called <code>exec</code>
ASU	process used superuser privileges (not in FreeBSD)
ACORE	process dumped core (not in Solaris10)
AXSIG	process was killed by a signal (not in Solaris10)
AEXPND	expanded accounting entry (Solaris10 only)
ANVER	new record format (FreeBSD only)

- The data required for the accounting record are all kept by the kernel in the process table and initialized whenever a new process is created.
- Each accounting record is written when the process terminates.
- The accounting records correspond to processes (pid), not programs.
- Example, (Figure 8.28 `proc/test1.c`, Figure 8.29 `proc/pracct.c`).

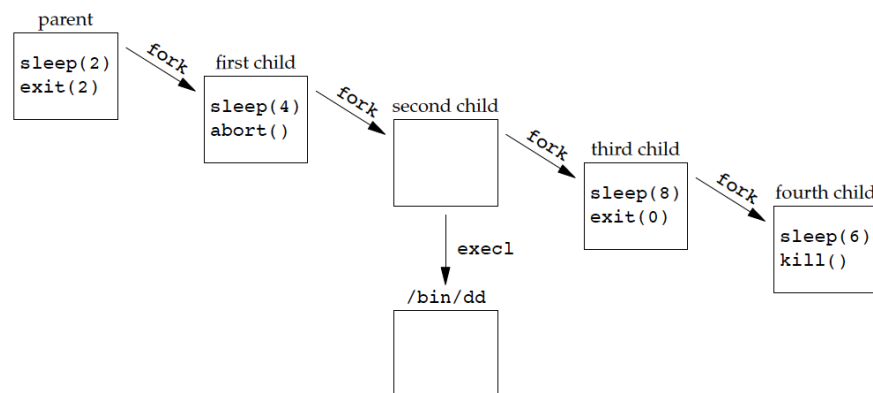


Figure 4: Process structure for accounting example

14 User Identification

- Sometime we want to find out the login name of the user who's running the program. Calling `getpwuid(getuid())` is a approach, but what if a single user has multiple login names?
- The system normally keeps track of the name we log in under, and the `getlogin` function provides a way to fetch that login name.

```
1 #include <unistd.h>
2 char *getlogin(void);
```

- This function fails if the process is not attached to a terminal that a user logged into.

15 Process Scheduling

- Historically, the UNIX System provided processes with only coarse control over their scheduling priority.
- The real-time extensions in POSIX added interfaces to select among multiple scheduling classes and fine-tune their behavior.
- In the Single UNIX Specification, nice values range from 0 to $(2 * \text{NZERO}) - 1$. Lower nice values have higher scheduling priority.
- A process can retrieve and change its own nice value with the `nice` function.

```
1 #include <unistd.h>
2 int nice(int incr);
```

- The `incr` argument is added to the nice value of the calling process. If `incr` is too large, the system silently reduces it to the maximum legal value. Similarly, if `incr` is too small, the system silently increases it to the minimum legal value.
- It returns `new nice value - NZERO` if OK, `-1` on error.
- Because `-1` is a legal successful return value, we need to clear `errno` before calling `nice` and check its value if `nice` returns `-1`.
- Example (Figure 8.30, `proc/nice.c`).
- The `getpriority` function can be used to get the nice value for a process, just like the `nice` function. However, `getpriority` can also get the nice value for a group of related processes.

```
1 #include <sys/resource.h>
2 int getpriority(int which, id_t who);
```

- It returns nice value between `-NZERO` and `NZERO-1` if OK, `-1` on error.
- The `which` argument can take on one of three values: `PRIO_PROCESS` to indicate a process, `PRIO_PGRP` to indicate a process group, and `PRIO_USER` to indicate a user ID.
- The `which` argument controls how the `who` argument is interpreted. If the `who` argument is 0, then it indicates the calling process, process group, or user (depending on the value of the `which` argument).
- When the `which` argument applies to more than one process, the highest priority (lowest value) of all the applicable processes is returned.

- The `setpriority` function can be used to set the priority of a process, a process group, or all the processes belonging to a particular user ID.

```
1 #include <sys/resource.h>
2 int setpriority(int which, id_t who, int value);
```

- The *which* and *who* arguments are the same as in the `getpriority` function. The *value* is added to `NZERO` and this becomes the new nice value.
- The SUS leaves it up to the implementation whether the nice value is inherited by a child process after a `fork`.

16 Process Times

- We knew that the `time` command can give a report of CPU usage of a give command.
- A process can call the `times` function to obtain these values for itself and any terminated children:

```
1 #include <sys/times.h>
2 clock_t times(struct tms *buf);
```

- The function fill in the `tms` structure pointed to by *buf*:

```
1 struct tms
2 {
3     clock_t tms_utime; /* User CPU time. */
4     clock_t tms_stime; /* System CPU time. */
5     clock_t tms_cutime; /* User CPU time of dead children. */
6     clock_t tms_cstime; /* System CPU time of dead children. */
7 };
```

- Note that the structure does not contain any measurement for the wall clock time. Instead, the function returns the wall clock time as the value of the function, each time it's called.
- Example (Figure 8.31, `proc/times1.c`).

The End of Chapter 8.