

Advanced Programming in the UNIX Environment — *Process Relationships*

Hop Lee
hoplee@bupt.edu.cn

SCHOOL OF INFORMATION AND COMMUNICATION ENGINEERING



Table of Contents

Terminal Login

Network Login

Process Groups

Sessions

Controlling Terminal

TCGETPGRP and TCSETPGRP Functions

Job Control

Shell Execution of Programs

Orphaned Process Groups



Terminal Login I

- ▶ The procedure that we describe is used to log in to a UNIX system using an terminal.
- ▶ When the system is bootstrapped the kernel creates process ID 1, the `init` process, and it is `init` that brings the system up multiuser.
- ▶ For every terminal device that allows a login in the file `/etc/ttys`, `init` does a `fork` followed by an `exec` of the program `getty`.



Terminal Login II




Figure: Process invoked by init to allow terminal logins.

- All the processes shown in Figure 1 have a real user ID 0 and an effective user ID 0.



Terminal Login III

- ▶ The terminal is opened for reading and writing by `getty` calling `open` for the terminal device.
- ▶ Once the device is opened, file descriptors 0, 1, and 2 are set to the device.
- ▶ Then `getty` outputs some common information about the system (`/etc/issue`) and a string like `login:` and waits for us to enter our user name.
- ▶ `getty` is done when we enter our user name. It then invokes the `login` program, similar to

```
1 execl("/usr/bin/login", "login", "-p",  
2   username, (char *)0, envp);
```



Terminal Login IV

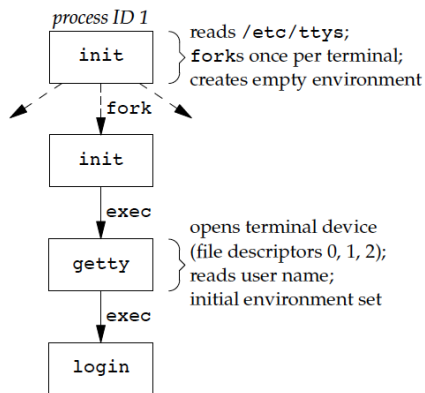


Figure: State of processes after login has been invoked



Terminal Login V

- ▶ All the processes shown in Figure 2 have superuser privileges.
- ▶ The process ID of the bottom three processes is the same.
- ▶ The `login` can call `getpwnam` to fetch our password file entry.
- ▶ Then calls `getpass` to display the prompt Password: and read our password.
- ▶ It calls `crypt` to encrypt the password that we entered and compares the result with the `pw_passwd` field.
- ▶ If the login attempt fails (after a few tries), `login` calls `exit` with an argument of 1. This termination will be noticed by the parent (`init`) and it will do another `fork` followed by an `exec` of `getty` again.



Terminal Login VI

- ▶ If we login correctly, `login` changes to our home directory (`chdir`), the ownership of our terminal device is changed (`chown`), the access permissions are also changed for the terminal device, our group IDs are set (`setgid` and then `initgroups`), the environment is initialized.
- ▶ Finally it changes to our user ID (`setuid`) and invokes our login shell, as in

```
1 execl("/bin/sh", "-sh", (char *)0);
```

- ▶ `login` really does more than we've described here. It optionally prints the message-of-the-day file (`/etc/motd`), checks for new mail, and does other functions.



Terminal Login VII

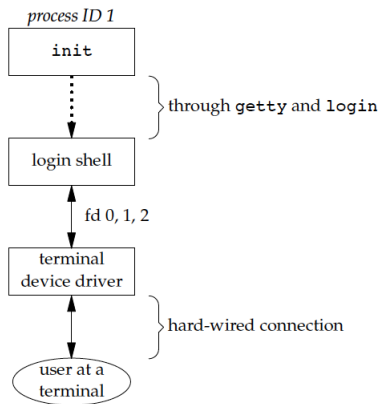


Figure: Arrangement of processes after everything is set for a terminal login



Terminal Login VIII

- ▶ Our login shell now reads its start-up files. These start-up file usually change come of the environment variables and add many additional variables to the environment.



Network Login I

- ▶ With the terminal logins, `init` knows which terminal devices are enabled for logins and spawns a `getty` process for each device.
- ▶ In the case of network logins, however, all the logins come through the kernel's network interface drivers, and we do not know ahead of time how many of these will occur.
- ▶ And there isn't any terminal ready for network logins.
- ▶ To resolve the problem, `init` invokes a shell to execute a script, and then create a child process named `inetd` (or `xinetd`).
- ▶ `inetd` waits for TCP/IP connection requests. When such a request arrives the host, it does a `fork` and `exec` of the appropriate program, eg. `telnetd`.



Network Login II

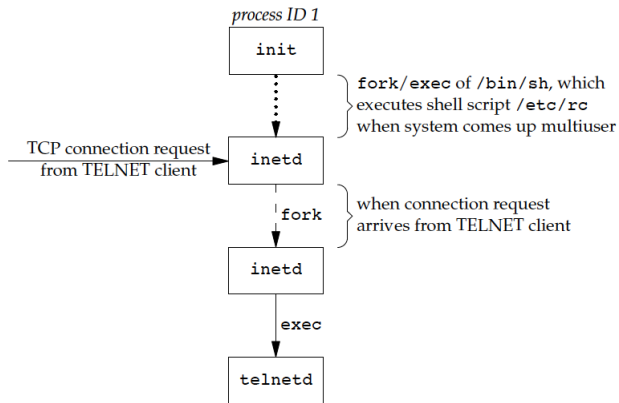


Figure: Sequence of processes involved in executing TELNET server



Network Login III

- ▶ The `telnetd` process then opens a pseudo-terminal device and splits into two processes using `fork`. The parent handles the communication across the network connection, and the child does an `exec` of the `login` program.



Network Login IV

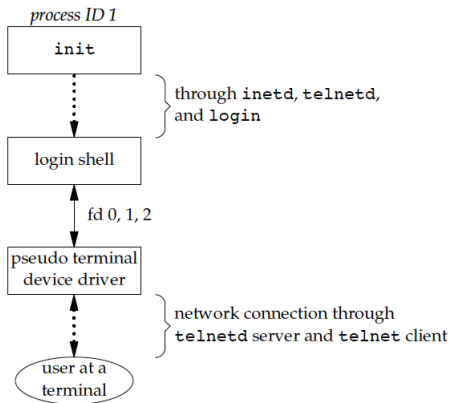


Figure: Arrangement of processes after everything is set for a network login



Network Login V

- ▶ All the procedures we just talked about is based on FreeBSD, and the other three experimental system are almost the same as it.



Process Groups I

- ▶ In addition to having a process ID, each process also belongs to a **process group**.
- ▶ A process group is a collection of one or more processes. Each process group has a unique process group ID. The function `getpgrp` returns the process group ID of the calling process, and the `getpgid` function took a *pid* argument and returned the process group for that process.

```
1 #include <unistd.h>
2 pid_t getpgrp(void);
3 pid_t getpgid(pid_t pid);
```



Process Groups II

- ▶ Each process group have a **process group leader**. The leader is identified by having its process group ID equal its process ID.
- ▶ The process group still exists, as long as there is at least one process in the group, regardless whether the group leader terminates or not.
- ▶ This is called the **process group lifetime** — the period of time that begins when the group is created and ends when the last remaining process in the group either terminate or enter some other process group.
- ▶ A process joins an existing process group, or creates a new process group by calling `setpgid`



Process Groups III

```
1 #include <unistd.h>
2 int setpgid(pid_t pid, pid_t pgid);
```

- ▶ This sets the process group ID to *pgid* of the process *pid*. If the two arguments are equal, the process specified by *pid* becomes a process group leader.
- ▶ A process can set the process group ID of only itself or one of its children.
- ▶ If *pgid* is 0, the process ID of the caller is used. Also, if *pgid* is 0, the process ID specified by *pid* is used as the process group ID.



Sessions I

- ▶ A **session** is a collection of one or more process groups.

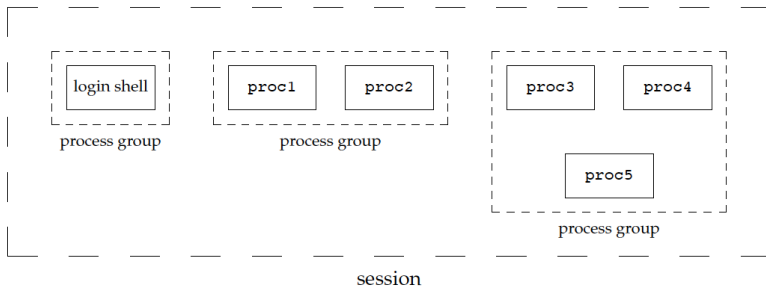


Figure: Arrangement of processes into process groups and sessions



Sessions II

- ▶ The processes in a process group are usually grouped together into the process group by a shell pipeline. For example, the arrangement shown in Figure 9.6 could have been generated by shell commands of the form

```
1 proc1 | proc2 &  
2 proc3 | proc4 | proc5
```

- ▶ A process establishes a new session by calling the `setsid` function.

```
1 #include <unistd.h>  
2 pid_t setsid(void);
```



Sessions III

- ▶ This function returns an error if the caller is already a process group leader. To ensure this is not the case, the usual practice is to call `fork` and have the parent terminate and the child continue.
- ▶ If the calling process is **NOT** a process group leader, this function creates a new session. Three things happen:
 1. The process becomes the **session leader** of this new session.
 2. The process becomes the **process group leader** of a new process group.
 3. The process has no **controlling terminal**. If the process had a controlling terminal before calling `setsid`, that association is broken.



Sessions IV

- ▶ SUS talks about a session ID that is the process ID of the **session leader**. The `getsid` function returns the process group ID of a process's session leader.

```
1 #include <unistd.h>
2 pid_t getsid(pid_t pid);
```

- ▶ It returns: session leader's process group ID if OK, `-1` on error
- ▶ If `pid` is 0, `getsid` returns the process group ID of the calling process's session leader.
- ▶ For security reasons, some implementations may restrict the calling process from obtaining the process group ID of the session leader if `pid` doesn't belong to the same session as the caller.



Controlling Terminal I

- ▶ A session can have a single **controlling terminal**. This is usually the terminal device or pseudo-terminal device on which we log in.
- ▶ The session leader that establishes the connection to the controlling terminal is called the **controlling process**.
- ▶ The process groups within a session can be divided into a single **foreground process group** and one or more **background process groups**.
- ▶ If a session has a controlling terminal, then it has a single foreground process group, and all other process groups in the session are background process groups.



Controlling Terminal II

- ▶ Whenever we type our terminal's interrupt key or quit key, this causes either the interrupt or quit signal to be sent to all processes in the foreground process group.
- ▶ If a modem/network disconnect is detected by the terminal interface, the hang-up signal is sent to the controlling process.



Controlling Terminal III

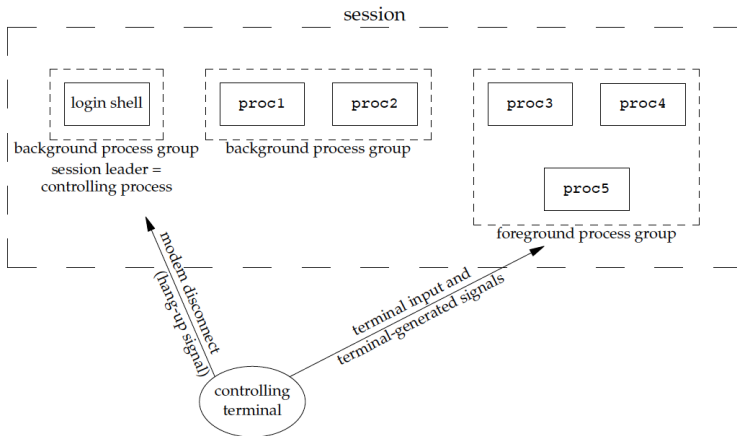


Figure: Arrangement of processes into process groups and sessions

Controlling Terminal IV

- ▶ Usually, we don't have to worry about the controlling terminal; it is established automatically when we log in.
- ▶ The way a program guarantees that it is talking to the controlling terminal is to open `/dev/tty`.



tcgetpgrp and tcsetpgrp Functions I

- ▶ The `tcgetpgrp` function can tell the kernel which process group is the foreground process group, so that the terminal device driver knows where to send the terminal input and the terminal-generated signals.

```
1 #include <unistd.h>
2 pid_t tcgetpgrp(int fd);
3 int tcsetpgrp(int fd, pid_t pgrp);
```

- ▶ The function `tcgetpgrp` returns the process group ID of the foreground process group associated with the terminal open on `fd`.



tcgetpgrp and tcsetpgrp Functions II

- ▶ If the process has a controlling terminal, the process can call `tcsetpgrp` to set the foreground process group ID to *pgrp*. The value of *pgrp* must be the process group ID of a process group in the same session, and *fd* must refer to the controlling terminal of the session.
- ▶ Most applications don't call these two functions directly. Instead, the functions are normally called by job-control shells.
- ▶ The `tcgetsid` function allows an application to obtain the process group ID for the session leader given a file descriptor for the controlling TTY.

```
1 #include <termios.h>
2 pid_t tcgetsid(int fd);
```



tcgetpgrp and tcsetpgrp Functions III

- ▶ It returns session leader's process group ID if OK, -1 on error.
- ▶ Applications that need to manage controlling terminals can use `tcgetsid` to identify the session ID of the controlling terminal's session leader.



Job Control I

- ▶ Job Control means we can start multiple jobs (groups of processes) from a single terminal and control which jobs can access the terminal and which jobs are run in the background.
- ▶ Job control requires three forms of support:
 1. A shell that supports job control
 2. The terminal driver in the kernel must support job control
 3. The kernel must support certain job-control signals
- ▶ Not all hosts support job control and not all shells support job control.
- ▶ The command **bg**, **fg**, **jobs**, and suspend character were prepared for job control under **bash**.
- ▶ The terminal driver looks for three special characters, which generate signals to the foreground process group.



Job Control II

- ▶ The interrupt character (typically DELETE or Control-C) generates SIGINT.
 - ▶ The quit character (typically Control-backslash) generates SIGQUIT.
 - ▶ The suspend character (typically Control-Z) generates SIGTSTP.
-
- ▶ When we enter characters at a terminal, only the foreground job receives terminal input.
 - ▶ It is not an error for a background job to try to read from the terminal, but the terminal driver detects this and sends a special signal to the background job: **SIGTIN**.
 - ▶ This signal normally stops the background job; by using the shell, we are notified of this event and can bring the job into the foreground so that it can read from the terminal.



Job Control III

- ▶ When a background job sends its output to the controlling terminal, there is an option that we can allow or disallow. Normally, we use the `stty(1)` command to change this option.
- ▶ Check Figure 9.9 in textbook for summary of the features of job control.



Shell Execution of Programs

- ▶ A long (complicated) example on textbook, p303-p307, shows how the shells execute programs and how this relates to the concepts of process groups, controlling terminals, and sessions.



Orphaned Process Groups I

- ▶ The POSIX.1 definition of an **orphaned process group** is one in which the parent of every member is either itself a member of the group or is not a member of the group's session.
- ▶ If the process group is not orphaned, there is a chance that one of those parents in a different process group but in the same session will restart a stopped process in the process group that is not orphaned.
- ▶ Since the process group is orphaned when the parent terminates, and the process group contains a stopped process, POSIX.1 requires that every process in the newly orphaned process group be sent the hang-up signal (**SIGHUP**) followed by the continue signal (**SIGCONT**).



Orphaned Process Groups II

- ▶ This causes the child to be continued, after processing the hang-up signal. The default action for the hang-up signal is to terminate the process, so we have to provide a signal handler to catch the signal.
- ▶ When a process in a background process group tries to read from its controlling terminal, `SIGTTIN` is generated for the background process group. But here we have an orphaned process group; if the kernel were to stop it with this signal, the processes in the process group would probably never be continued. POSIX.1 specifies that the `read` is to return an error with `errno` set to `EIO` in this situation.
- ▶ The child was placed in a background process group when the parent terminated.



The End

The End of Chapter 9.

